

Kurs Bio144:

Datenanalyse in der Biologie

Stefanie Muff & Owen L. Petchey

Lecture 7: ANCOVA, short introduction to Linear Algebra

1. April 2019

Overview

- ANCOVA
- Introduction to linear Algebra

Note:

ANCOVA = ANalysi of COVAriance (Kovarianzanalyse)

Course material covered today

The lecture material of today is based on the following literature:

- “Getting Started with R” chapter 6.3
- “Lineare regression” chapters 3.A (p. 43-45) and 3.4, 3.5 (p. 39-42)


Recap of ANOVA

- ANOVA is a method to test if the means of **two or more groups are different**.
- Post-hoc tests and contrasts, including correction for p -values, to understand the differences between the groups.
- Two-way ANOVA for factorial designs, interactions.
- ANOVA is a special case of linear regression with categorical covariates.

Recap of two-way ANOVA example

Remember: Influence of four levels of fertilizer (DUENGER) on the yield (ERTRAG) on 5 species (SORTE) of crops was investigated. For each DUENGER \times ERTRAG combination, 3 repeats were taken.

Interaction plot with ERTRAG and $\log(\text{ERTRAG})$ as response:



Remember: We used $\log(\text{ERTRAG})$, because the residual plots were otherwise not ok.

```

r.duenger2 <- lm(log(ERTRAG) ~ DUENGER*SORTE,d.duenger)
anova(r.duenger2)

## Analysis of Variance Table
##
## Response: log(ERTRAG)
##           Df  Sum Sq Mean Sq  F value Pr(>F)
## DUENGER      3 11.6917   3.8972 854.0505 <2e-16 ***
## SORTE        4   8.5202   2.1300 466.7851 <2e-16 ***
## DUENGER:SORTE 12   0.0929   0.0077   1.6958 0.1045
## Residuals    40   0.1825   0.0046
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Questions:

- Number of parameters?
- Degrees of freedom (60 data points)?
- Interpretation?

```
summary(r.duenger2)

##
## Call:
## lm(formula = log(ERTRAG) ~ DUENGER * SORTE, data = d.duenger)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.120968 -0.045595  0.008984  0.049072  0.102175
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.68505    0.03900   68.846 < 2e-16 ***
## DUENGER2        0.43165    0.05516    7.826 1.36e-09 ***
## DUENGER3        0.79997    0.05516   14.504 < 2e-16 ***
## DUENGER4        1.21152    0.05516   21.966 < 2e-16 ***
## SORTE2          0.38979    0.05516    7.067 1.51e-08 ***
## SORTE3          0.55799    0.05516   10.117 1.38e-12 ***
## SORTE4          0.82018    0.05516   14.870 < 2e-16 ***
## SORTE5          1.08169    0.05516   19.612 < 2e-16 ***
## DUENGER2:SORTE2 -0.12949    0.07800   -1.660  0.105
## DUENGER3:SORTE2 -0.10613    0.07800   -1.361  0.181
## DUENGER4:SORTE2 -0.04924    0.07800   -0.631  0.531
## DUENGER2:SORTE3 -0.12180    0.07800   -1.562  0.126
## DUENGER3:SORTE3 -0.18034    0.07800   -2.312  0.026 *
## DUENGER4:SORTE3 -0.16061    0.07800   -2.059  0.046 *
## DUENGER2:SORTE4 -0.10138    0.07800   -1.300  0.201
## DUENGER3:SORTE4 -0.05311    0.07800   -0.681  0.500
## DUENGER4:SORTE4 -0.02954    0.07800   -0.379  0.707
## DUENGER2:SORTE5 -0.08779    0.07800   -1.125  0.267
## DUENGER3:SORTE5  0.04370    0.07800    0.560  0.578
## DUENGER4:SORTE5  0.09014    0.07800    1.156  0.255
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06755 on 40 degrees of freedom
## Multiple R-squared:  0.9911, Adjusted R-squared:  0.9869
```

Analysis of Covariance (ANCOVA)

An ANCOVA is an analysis of variance (ANOVA), **including** also at least one **continuous covariate**.

ANCOVA unifies several concepts that we approached in this course so far:

- Linear regression
- Categorical covariates
- Interactions (of continuous and categorical covariates)
- Analysis of Variance (ANOVA)

As such, it is a **special case of the linear regression model**.

Given a categorical covariate x_i and a continuous covariate z_i . Then the ANCOVA equation (without interactions) is given as

$$y_i = \beta_0 + \beta_1 x_i^{(1)} + \dots + \beta_k x_i^{(k)} + \beta_z z_i + \epsilon_i ,$$

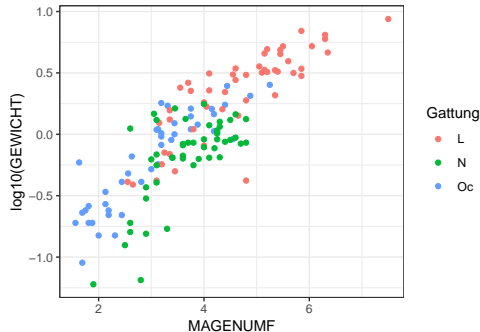
where $x_i^{(k)}$ is the k th dummy variable ($x_i^{(k)}=1$ if i th observation belongs to category k , 0 otherwise).

Note 1: It is straightforward to add an interaction of x_i with z_i .

Note 2: Again, for identifiability reason, we typically set $\beta_1 = 0$.

Once more: the earthworms

“Magenumfang” was used to predict “Gewicht” of the worm, including as covariate also the worm species.



Categorical and continuous covariates were used to predict a continuous outcome → ANCOVA.

```
r.lm <- lm(log(GEWICHT) ~ MAGENUMF + Gattung,d.wurm)
summary(r.lm)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	-2.5355459	0.22147279	-11.4485663	8.617670e-22
## MAGENUMF	0.7118725	0.04528843	15.7186392	1.232126e-32
## GattungN	-0.5151344	0.11009219	-4.6791186	6.760621e-06
## GattungOc	-0.0907298	0.12791000	-0.7093254	4.793107e-01

Important: The p -values for the entries `GattungN` and `GattungOc` are not very meaningful (why?).

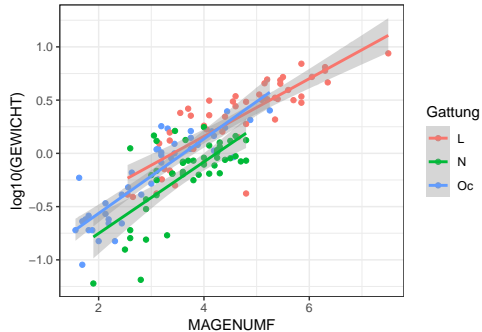
To understand if “Gattung” has an effect, **we need to carry out an F -test**
→ ANOVA table:

```
anova(r.lm)
```

```
## Analysis of Variance Table
##
## Response: log(GEWICHT)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
## MAGENUMF	1	104.866	104.866	409.69	< 2.2e-16 ***
## Gattung	2	7.177	3.589	14.02	2.842e-06 ***
## Residuals	139	35.579	0.256		
## ---					
## Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' ' 1

We also included an **interaction term** between MAGENUMF and Gattung to allow for different slopes:



→ We again need the *F*-test to check whether the respective interaction term is needed:

```
r.lm2<- lm(log(GEWICHT) ~ MAGENUMF * Gattung,d.wurm)
anova(r.lm2)

## Analysis of Variance Table
##
## Response: log(GEWICHT)
##              Df Sum Sq Mean Sq F value    Pr(>F)
## MAGENUMF      1 104.866  104.866 414.4743 < 2.2e-16 ***
## Gattung       2   7.177    3.589  14.1835 2.521e-06 ***
## MAGENUMF:Gattung 2   0.917    0.458   1.8112  0.1673
## Residuals    137  34.662    0.253
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

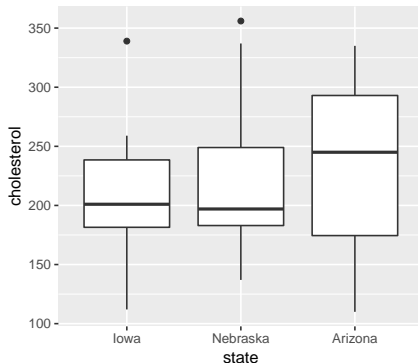
→ $p = 0.167$, thus interaction is probably not relevant.

A new example: cholesterol levels

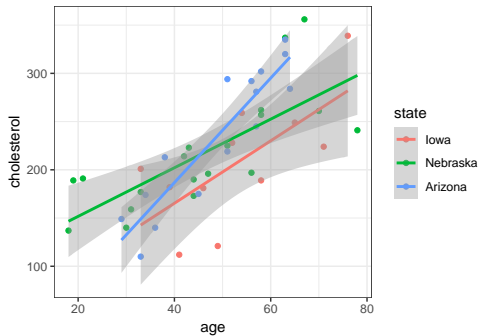
Example: Cholesterol levels [mg/ml] for 45 women from three US states (Iowa, Nebraska, Arizona), were measured.

Question: Do these levels differ between the states?

Age (years) may be a relevant covariable.



The scatter plot gives an idea about the model that might be useful here:



→ We include state, age and the interaction of the two.

Doing the analysis:

```
r.lm <- lm(cholesterol ~ age*state,data=d.chol)
anova(r.lm)

## Analysis of Variance Table
##
## Response: cholesterol
##           Df Sum Sq Mean Sq F value    Pr(>F)
## age         1  96524   96524  61.8961 1.424e-09 ***
## state        2  11474    5737   3.6789  0.03438 *
## age:state    2  12665    6332   4.0606  0.02501 *
## Residuals   39  60819    1559
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Interpretation?

Compare the results from the previous slide to the estimated coefficients:

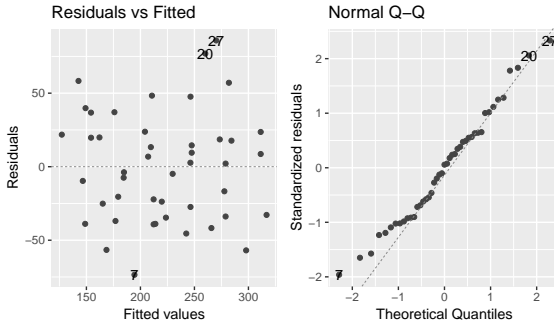
```
r.lm <- lm(cholesterol ~ age*state,data=d.chol)
summary(r.lm)$coef
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	35.8112138	50.4494150	0.7098440	0.482023781
## age	3.2381449	0.9234015	3.5067573	0.001158119
## stateNebraska	65.4865523	56.7347108	1.1542590	0.255418196
## stateArizona	-65.7063829	66.7677031	-0.9841043	0.331130339
## age:stateNebraska	-0.7177069	1.0643772	-0.6742975	0.504099737
## age:stateArizona	2.1787633	1.2672928	1.7192264	0.093502039

Note: The p -values for the age coefficient is not the same as in the ANOVA table.

Reason: `anova()` tests the models against one another in the **order** specified.

As always, some model checking is necessary:



→ This seems ok.

An introduction to linear Algebra

Who has some knowledge of linear Algebra?

Overview

- The basics about
 - vectors
 - matrices
 - matrix algebra
 - matrix multiplication
- Why is linear Algebra useful?
- What does it have to do with data analysis and statistics?
- Regression equations in matrix notation.

Motivation

Why are vectors, matrices and their algebraic rules useful?

Example 1: The observations for a covariate \mathbf{x} or the response \mathbf{y} for all individuals $1 \leq i \leq n$ can be stored in a vector (vectors and matrices are always given in **bold** letters):

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}, \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{pmatrix}.$$

Example 2: Covariance matrices for multiple variables. Say we have $\mathbf{x}^{(1)}$ and $\mathbf{x}^{(2)}$. The **covariance matrix** is then given as

$$\begin{pmatrix} \text{Var}(\mathbf{x}^{(1)}) & \text{Cov}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) \\ \text{Cov}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) & \text{Var}(\mathbf{x}^{(2)}) \end{pmatrix}.$$

Example 3: The **data** (e.g. of some regression model) can be stored in a **matrix**:

$$\tilde{\mathbf{X}} = \begin{pmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_2^{(1)} & x_2^{(2)} \\ \dots & \dots & \dots \\ 1 & x_n^{(1)} & x_n^{(2)} \end{pmatrix} .$$

This is the so-called **design matrix** with a vector of 1's in the first column.

Example 4: A linear regression model can be written compactly using **matrix multiplication**:

$$\mathbf{y} = \tilde{\mathbf{X}} \cdot \tilde{\boldsymbol{\beta}} + \mathbf{e} ,$$

with $\tilde{\boldsymbol{\beta}}$ the vector of regression coefficients and \mathbf{e} the vector of errors

Why do we discuss this topic in our course?

- Useful for **compact notation**.
- Enables you to **understand many statistical texts** (books, research articles) that remain inaccessible otherwise.
- Useful for **efficient coding**, e.g. in R, which helps to increase speed and to reduce error rates.
- More advanced concepts often rely on linear algebra, e.g. **principal component analysis** (PCA) or **random effects** models.
- Is part of a **general education** (Allgemeinbildung) ;-)

Matrices

An $n \times m$ **Matrix** is given as

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1m} \\ a_{21} & a_{22} & \dots & a_{2m} \\ \vdots & \vdots & & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nm} \end{pmatrix},$$

with rows $i = 1, \dots, n$ and columns $j = 1, \dots, m$.

Quadratic matrix: $n = m$. Example:

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 3 & 2 \\ 6 & 1 & 9 \end{pmatrix}$$

Symmetric matrix: $a_{ij} = a_{ji}$. Example:

$$\begin{pmatrix} 1 & 2 & 3 \\ 2 & 3 & 4 \\ 3 & 4 & 5 \end{pmatrix} .$$

The diagonal of a quadratic matrix is given by $(a_{11}, a_{22}, \dots, a_{nn})$.

Example: the diagonal of the above matrix is given as

$$(a_{11}, a_{22}, a_{33}) = (1, 3, 5) .$$

Diagonal matrix: A matrix that has entries $\neq 0$ **only on the diagonal**.

Example:

$$\begin{pmatrix} 1 & 0 & 0 \\ 0 & 3 & 0 \\ 0 & 0 & 5 \end{pmatrix} .$$

Transposing a matrix: Given a matrix **A**. Exchange the rows by the columns and vice versa. This leads to the **transposed matrix** \mathbf{A}^\top :

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1m} \\ a_{21} & a_{22} & \dots & a_{2m} \\ \vdots & \vdots & & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nm} \end{pmatrix} \Rightarrow \mathbf{A}^\top = \begin{pmatrix} a_{11} & a_{21} & \dots & a_{n1} \\ a_{12} & a_{22} & \dots & a_{n2} \\ \vdots & \vdots & & \vdots \\ a_{1m} & a_{2m} & \dots & a_{nm} \end{pmatrix}.$$

Examples (note also the change of dimensions):

$$\mathbf{A} = \begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 13 & 14 & 15 & 16 \end{pmatrix} \Rightarrow \mathbf{A}^\top = \begin{pmatrix} 1 & 5 & 9 & 13 \\ 2 & 6 & 10 & 14 \\ 3 & 7 & 11 & 15 \\ 4 & 8 & 12 & 16 \end{pmatrix}$$

$$\mathbf{A} = \begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \end{pmatrix} \Rightarrow \mathbf{A}^\top = \begin{pmatrix} 1 & 5 \\ 2 & 6 \\ 3 & 7 \\ 4 & 8 \end{pmatrix}$$

- Transposing a matrix **twice** leads to the original matrix:

$$(\mathbf{A}^\top)^\top = \mathbf{A} .$$

- When a matrix is **symmetric**, then

$$\mathbf{A}^\top = \mathbf{A} .$$

This is true in particular for diagonal matrices.

Vectors

A vector is nothing else than n numbers written in a column:

$$\mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix}$$

Transposing a vector leads to a *row vector*:

$$\begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix}^{\top} = \begin{pmatrix} b_1 & b_2 & \dots & b_n \end{pmatrix}$$

Note: By definition (by default), a vector is always a column vector.

Addition and subtraction

- Adding and subtracting matrices and vectors is only possible when the objects have the **same dimension**.
- Examples: Elementwise addition (or subtraction)

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix} + \begin{pmatrix} 3 & 2 & 1 \\ 6 & 5 & 4 \end{pmatrix} = \begin{pmatrix} 4 & 4 & 4 \\ 10 & 10 & 10 \end{pmatrix}$$

$$\begin{pmatrix} 1 \\ 4 \end{pmatrix} - \begin{pmatrix} 3 \\ 9 \end{pmatrix} = \begin{pmatrix} -2 \\ -5 \end{pmatrix}$$

- But this addition is **not defined**:

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix} + \begin{pmatrix} 3 & 6 \\ 2 & 5 \\ 1 & 4 \end{pmatrix} =$$

Multiplication by a scalar

Multiplication with a “number” (scalar) is simple: Multiply each element in a vector or a matrix.

Examples:

$$3 \cdot \begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix} = \begin{pmatrix} 3 & 6 & 9 \\ 12 & 15 & 18 \end{pmatrix}$$

$$-2 \cdot \begin{pmatrix} 1 \\ 4 \\ -2 \end{pmatrix} = \begin{pmatrix} -2 \\ -8 \\ 4 \end{pmatrix}$$

Matrix multiplication

The multiplication of two matrices **A** and **B** is **defined if**
number of columns in **A** = number of rows in **B**.

It is easiest to explain matrix multiplication with an example:

$$\begin{pmatrix} 2 & 1 \\ -1 & 0 \\ 3 & 1 \end{pmatrix} \cdot \begin{pmatrix} 3 & 1 \\ 4 & -2 \end{pmatrix} = \begin{pmatrix} 2 \cdot 3 + 1 \cdot 4 & 2 \cdot 1 + 1 \cdot (-2) \\ (-1) \cdot 3 + 0 \cdot 4 & (-1) \cdot 1 + 0 \cdot (-2) \\ 3 \cdot 3 + 1 \cdot 4 & 3 \cdot 1 + 1 \cdot (-2) \end{pmatrix}$$
$$= \begin{pmatrix} 10 & 0 \\ -3 & -1 \\ 13 & 1 \end{pmatrix}$$

► Matrix multiplication app

Matrix multiplication rules I

Attention: Matrix multiplication does **not** follow the same rules as scalar multiplication!!!

- It can happen that $\mathbf{A} \cdot \mathbf{B}$ can be calculated, but $\mathbf{B} \cdot \mathbf{A}$ is not defined (see example on previous slide).
- In general: $\mathbf{A} \cdot \mathbf{B} \neq \mathbf{B} \cdot \mathbf{A}$, even if both are defined.
- It can happen that $\mathbf{A} \cdot \mathbf{B} = \mathbf{0}$ (0 matrix), although both $\mathbf{A} \neq \mathbf{0}$ and $\mathbf{B} \neq \mathbf{0}$.
- The **Assoziativgesetz** holds: $\mathbf{A} \cdot (\mathbf{B} \cdot \mathbf{C}) = (\mathbf{A} \cdot \mathbf{B}) \cdot \mathbf{C}$.
- The **Distributivgesetz** holds:

$$\mathbf{A} \cdot (\mathbf{B} + \mathbf{C}) = \mathbf{A} \cdot \mathbf{B} + \mathbf{A} \cdot \mathbf{C}$$

$$(\mathbf{A} + \mathbf{B}) \cdot \mathbf{C} = \mathbf{A} \cdot \mathbf{C} + \mathbf{B} \cdot \mathbf{C}$$

Matrix multiplication rules II

- Transposing inverts the order: $(\mathbf{A} \cdot \mathbf{B})^\top = \mathbf{B}^\top \cdot \mathbf{A}^\top$.
- The product $\mathbf{A} \cdot \mathbf{A}^\top$ is **always symmetric**.
- All these rules also hold for **vectors**, which can be interpreted as $n \times 1$ matrices:

$$\mathbf{a} \cdot \mathbf{b}^\top = \begin{pmatrix} a_1 b_1 & a_1 b_2 & \dots & a_1 b_m \\ a_2 b_1 & a_2 b_2 & \dots & a_2 b_m \\ \vdots & \vdots & & \vdots \\ a_n b_1 & a_n b_2 & \dots & a_n b_m \end{pmatrix}$$

If \mathbf{a} and \mathbf{b} have the **same length**:

$$\mathbf{a}^\top \cdot \mathbf{b} = \sum_i a_i b_i$$

Short exercises

Given vectors **a** and **b** and matrix **C**:

$$\mathbf{a} = \begin{pmatrix} 1 \\ -2 \\ 3 \\ 0 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} -2 \\ 4 \end{pmatrix}, \quad \mathbf{C} = \begin{pmatrix} 1 & 0 \\ -1 & 2 \end{pmatrix}$$

Calculate, if defined

- $\mathbf{a}^T \cdot \mathbf{b}$
- $\mathbf{a} \cdot \mathbf{b}^T$
- $\mathbf{C} \cdot \mathbf{a}$
- $\mathbf{C} \cdot \mathbf{b}$

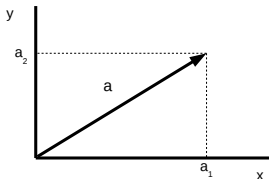
The length of a vector

The **length of a vector** $\mathbf{a}^T = (a_1, a_2, \dots, a_n)$ is defined as $\|\mathbf{a}\|$ with

$$\|\mathbf{a}\|^2 = \mathbf{a}^T \cdot \mathbf{a} = \sum_i a_i^2.$$

This is basically the **Pythagoras** idea in 2, 3, ... n dimensions.

In 2 dimensions: $\|\mathbf{a}\| = \sqrt{a_1^2 + a_2^2}$:



Identity matrix (Einheitsmatrix)

The identity matrix (of dimension m) is probably the simplest matrix that exists. It has 1's on the diagonal and 0's everywhere else:

$$\mathbf{I} = \begin{pmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & & \vdots \\ 0 & 0 & \dots & 1 \end{pmatrix}$$

Multiplication with the identity matrix leaves a $m \times n$ matrix \mathbf{A} unchanged:

$$\mathbf{A} \cdot \mathbf{I} = \mathbf{A} .$$

Inverse matrix

Given a quadratic matrix \mathbf{A} that fulfills

$$\mathbf{B} \cdot \mathbf{A} = \mathbf{I} ,$$

then \mathbf{B} is called the **inverse** of \mathbf{A} (and vice versa). One then writes

$$\mathbf{B} = \mathbf{A}^{-1} .$$

Note:

- In that case it also holds that $\mathbf{A} \cdot \mathbf{B} = \mathbf{I}$.
- Therefore: $\mathbf{A} = \mathbf{B}^{-1} \Leftrightarrow \mathbf{B} = \mathbf{A}^{-1}$

- The inverse of \mathbf{A} may **not exist**. If it exists, \mathbf{A} is **regular**, otherwise **singular**.
- $(\mathbf{A}^{-1})^{-1} = \mathbf{A}$.
- The inverse of a matrix product is given as

$$(\mathbf{A} \cdot \mathbf{B})^{-1} = \mathbf{B}^{-1} \cdot \mathbf{A}^{-1} .$$

- It is

$$(\mathbf{A}^{\top})^{-1} = (\mathbf{A}^{-1})^{\top} .$$

Therefore one may also write $\mathbf{A}^{-\top}$.

Linear regression in matrix notation

Linear regression with n data points can be understood as an **equation system with n equations**.

Remember example 4 from slide 21: We said that a linear regression model can be written compactly using **matrix multiplication**:

$$\mathbf{y} = \tilde{\mathbf{X}} \cdot \tilde{\boldsymbol{\beta}} + \mathbf{e} .$$

Task: Verify this now, using a model with two variables $\mathbf{x}^{(1)}$ and $\mathbf{x}^{(2)}$ and

$$\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}, \quad \tilde{\mathbf{X}} = \begin{pmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_2^{(1)} & x_2^{(2)} \\ \dots & \dots & \dots \\ 1 & x_n^{(1)} & x_n^{(2)} \end{pmatrix}, \quad \tilde{\boldsymbol{\beta}} = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{pmatrix}, \quad \mathbf{e} = \begin{pmatrix} e_1 \\ e_2 \\ \vdots \\ e_n \end{pmatrix} .$$

It can be shown (see Stahel 3.4f,g) that the least-squares estimates $\hat{\beta}$ can be calculated as

$$\hat{\beta} = (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \cdot \tilde{\mathbf{X}}^\top \cdot \mathbf{y}$$

Does this look complicated?

Let's test this in R

Doing linear algebra in R

Let us look at model $\mathbf{y} = \tilde{\mathbf{X}} \cdot \tilde{\boldsymbol{\beta}} + \mathbf{e}$ with coefficients $\beta_0 = 10, \beta_1 = 5, \beta_2 = -2$ and variables

i	$x_i^{(1)}$	$x_i^{(2)}$
1	0	4
2	1	1
3	2	0
4	3	1
5	4	4

Thus the model is given as

$$y_i = 10 + 5x_i^{(1)} - 2x_i^{(2)} + \epsilon_i, \text{ for } 1 \leq i \leq n.$$

Let us start by generating the “true” response, calculated as $\tilde{\mathbf{X}}\tilde{\boldsymbol{\beta}}$:


```

x1 <- c(0,1,2,3,4)
x2 <- c(4,1,0,1,4)
Xtilde <- matrix(c(rep(1,5),x1,x2),ncol=3)
Xtilde

##      [,1] [,2] [,3]
## [1,]    1    0    4
## [2,]    1    1    1
## [3,]    1    2    0
## [4,]    1    3    1
## [5,]    1    4    4

t.beta <- c(10,5,-2)
t.y <- Xtilde%*%t.beta
t.y

##      [,1]
## [1,]    2
## [2,]   13
## [3,]   20
## [4,]   23
## [5,]   22

```

Matrix multiplication in R is done by the **%*%** symbol.

Next, we generate the vector containing the $\epsilon_i \sim N(0, \sigma^2)$ with $\sigma^2 = 1$:

```
t.e <- rnorm(5,0,1)
t.e

## [1] 0.7606833 -0.3257157 0.6830309 0.9070262 0.9342162
```

which we add to the “true” $y = \tilde{X}\tilde{\beta}$ values, to obtain the “observed” values:

```
t.Y <- t.y + t.e
t.Y

##           [,1]
## [1,]  2.760683
## [2,] 12.674284
## [3,] 20.683031
## [4,] 23.907026
## [5,] 22.934216
```

It is now possible to fit the model with `lm`:

```
r.lm <- lm(t.Y ~ x1 + x2)
summary(r.lm)$coef

##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 10.069826  0.5556231  18.12348 0.003030672
## x1           5.157981  0.1866953  27.62780 0.001307540
## x2          -1.896970  0.1577864 -12.02239 0.006847617
```

Alternatively, we can use formula

$$\hat{\beta} = (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top \mathbf{y}$$

to find the parameter estimates:

```
solve(t(Xtilde) %*% Xtilde) %*% t(Xtilde) %*% t.Y  
  
##           [,1]  
## [1,] 10.069826  
## [2,]  5.157981  
## [3,] -1.896970
```

- `solve()` calculates the **inverse** (here the inverse of $\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}}$).
- `t()` gives the **transposed** (here of $\tilde{\mathbf{X}}^\top$).

Task: Do this calculation by yourself and verify for each step that the dimensions of the matrices and the vector are indeed fitting, so that this expression is defined.

Appendix

Some R commands for matrix algebra

Reading vectors and a matrices into R:

```
a <- c(1,2,3)
```

```
a
```

```
## [1] 1 2 3
```

```
A <- matrix(c(1,2,3,4,5,6),byrow=T,nrow=2)
```

```
B <- matrix(c(6,5,4,3,2,1),byrow=T,nrow=2)
```

```
A
```

```
##      [,1] [,2] [,3]
```

```
## [1,]    1    2    3
```

```
## [2,]    4    5    6
```

```
B
```

```
##      [,1] [,2] [,3]
```

```
## [1,]    6    5    4
```

```
## [2,]    3    2    1
```

Adding and subtracting:

```
A + B
```

```
##      [,1] [,2] [,3]  
## [1,]    7    7    7  
## [2,]    7    7    7
```

```
A - B
```

```
##      [,1] [,2] [,3]  
## [1,]   -5   -3   -1  
## [2,]    1    3    5
```

However, be careful, R sometimes does unreasonable things:

```
A + a
```

```
##      [,1] [,2] [,3]  
## [1,]    2    5    5  
## [2,]    6    6    9
```

What happened here??

Matrix multiplication:

```
C <- A %*% t(B)
C

##      [,1] [,2]
## [1,]   28  10
## [2,]   73  28

A%*%a

##      [,1]
## [1,]   14
## [2,]   32
```

Matrix inversion (possible for quadratic matrices only):

```
solve(C)

##      [,1]      [,2]
## [1,] 0.5185185 -0.1851852
## [2,] -1.3518519  0.5185185

C %*% solve(C)

##      [,1] [,2]
## [1,]    1    0
## [2,]    0    1
```

Why does solve(A) or solve(B) not work?